


N-body problem

- N-body simulations are a class of computational problems where we calculate the effect of a force between N objects (or bodies)
- The problem is to calculate the positions and movements of a number of bodies in space as time advances
 - the bodies are affected by all other bodies via gravitation
 - long-range interactions: each object affects all other objects
- Here we will only consider the two-dimensional case
 - extension to three dimensions is straight forward
- We will only consider gravitational force
 - attraction between two bodies
 - can also use the same methods for other types of forces, for instance electrostatic attraction or repulsion

Problem description

- We have N bodies in 2-dimensional space
 - bodies are treated as point masses
 - shape and size does not affect behaviour
- Each body is described by its
 - mass m
 - position $X = (x_x, x_y)$
 - velocity $V = (v_x, v_y)$
the rate of change in position over time, $V = dX/dt$
 - acceleration $A = (a_x, a_y)$
the rate of change in velocity over time, $A = dV/dt$
- As a body is affected by a force, its velocity changes
 - Newton's laws describe how bodies in space affect each other with gravitation



force ← - - - - - →
velocity → - - - - - →

Laws of gravitation


■ Newton's laws

- two bodies a and b with masses m_a and m_b
- positions are $X_a = (x_a, y_a)$ and $X_b = (x_b, y_b)$
- gravitational force on a caused by b

- $F_{a,b} = \frac{Gm_a m_b}{r^2} \frac{X_b - X_a}{r}$ where G is the gravitational constant, $G = 6.67259 \times 10^{-11}$

- and $r = \sqrt{(x_b - x_a)^2 + (y_b - y_a)^2}$ is the distance between a and b


- forces are vectors with x- and y-components $F = (f_x, f_y)$
- pairwise forces are symmetric but of opposite direction: $F_{a,b} = -F_{b,a}$



Laws of gravitation (cont.)

- The total force on a body a is the sum of the pairwise forces from all the other bodies

$$- \quad F_a = \sum_{j=0, j \neq a}^{N-1} F_{a,j} = G m_a \sum_{j=0, j \neq a}^{N-1} m_j \left[\frac{X_j - X_a}{r_{a,j}^3} \right]$$



- The force on a body affects its motion according to Newton's second law
 $F = m * a$
 - $a = F/m$ (acceleration is force divided by mass)
- Given the current position of a body (x_a, y_a) and the acceleration we can compute the velocity and position of the body in the next time step

Discrete solution


- We divide the time into short time intervals of length h (often denoted Δt)
 - starting from an initial state at time t_0 , we calculate the position and velocity for each body at times t_1, t_2, t_3, \dots
 - if the current timestep is t_i then the next timestep is $t_{i+1} = t_i + h$
- $O(N^2)$ algorithm
 - for each time interval
 - for each body b
 - for all other bodies c
 - calculate the force on b caused by c
 - calculate new velocity for b
 - calculate new position for b
- The time interval h has to be short enough to give an accurate solution
 - the amount of computation increases with a shorter interval

Discrete formulation

- For a body with mass m we compute the force F affecting it at time t
 - sum of the forces from all other bodies
- Then we compute the velocity and position of the body for the next timestep, i.e. at time $t+h$
 - new velocity is $V^{t+1} = V^t + A^t h$ where $A^t = \frac{F}{m}$
 - new position is $X^{t+1} = X^t + V^t h$
- When the bodies move to new positions, the forces change and the computation has to be repeated
- Called Euler's method


Problem with Euler's method

- In Euler's method, we assume that the acceleration is constant during the time interval
 - we use the acceleration at time t to calculate the new velocity at time $t+h$
- Acceleration is not constant during a time interval
 - when a body gets closer to another body, its acceleration increases
 - our approximation is based on the velocity at time t
- Have to use a small time step to get reliable results




The Leapfrog method

- In the Leapfrog method we use the *mid-point* of the interval to approximate the velocity during the interval
 - this is used to calculate the positions at the next point in time
- Velocities and positions are not updated at the same point of time
 - update positions at the *beginning* of each time step and velocities at the *middle* of the time steps
- Positions are updated at times $h, 2h, 3h, \dots$
 - $X_{i+1} = X_i + h \cdot V_{i+1/2}$ for $i=0, 1, 2, \dots$
- Velocities are updated at times $1/2h, t+1/2h, 2t+1/2h, 3+1/2h, \dots$
 - $V_{i+1/2} = V(t+1/2h)$



Properties of the Leapfrog method

- Using the velocity at the midpoint of the interval gives a better approximation than the velocity at the beginning (or end) of the interval
- No more complicated to implement than the Euler scheme
- The initial velocity at $V_{1/2}$ can for instance be calculated by Eulers method
- The leapfrog method is time reversible
 - starting from any state at time t_i we can calculate backwards in time to t_0
- The leapfrog method is a second order approximation
 - the approximations of the positions have an accuracy of $O(\Delta t^2)$



Implementation of the Leapfrog method

- A sequential implementation of the Leapfrog method is

```
set initial values of V to the velocities at time 0.5*deltat


for (t=0; t<maxt; t++) {           /* For each timestep */
    for (i=0; i<nbodies; i++) {    /* For each body */
        X[i] = X[i]+V[i]*deltat;    /* Update position */
        F = Force(i);              /* Calculate force on i */
        V[i] = V[i]+F*deltat/m;    /* Update velocity */
    }
}
```

- First update positions, then forces and acceleration and finally velocities
 - no need to keep positions and velocities from the previous timestep
 - first move the bodies based on their current velocities
 - then calculate the new forces based on the new positions
 - finally, calculate the new velocities

Parallel N -body simulation: decomposition

■ N bodies, P processes

- we divide the bodies evenly among the processes
- $M = N/P$ local bodies per process




- A process updates the positions, forces and velocities *for its own bodies*
 - to calculate the forces, it needs information about the positions of *all the other* bodies
- The algorithm is the same as in a sequential solution, but it is only applied to the *own subset of bodies*
 - move the bodies, exchange updated positions with all other processes
 - calculate the forces on own bodies
 - calculate the new velocities for the own bodies

Parallel N -body simulation: updating the bodies

- Each process stores and updates its own range of the particles
 - each process has $M = N/P$ local bodies
 - can use variables *start* and *end* to store the interval of bodies for each process
 - $first = N*id/np$
 - $last = N*(id+1)/np$

where N is total number of bodies, id is the process rank and np is the number of processes



- Move the bodies to new positions:
 - similar code to update velocities

```
for (int i=first; i<last; i++){  
    X[i] = X[i] + Vx[i]*dt;  
    Y[i] = Y[i] + Vy[i]*dt;  
}
```

- Calculate forces for the updated positions:

```
ComputeForce_parallel(first, last, N, X, Y, mass, Fx, Fy);
```

Communication structures


- A parallel implementation of the N-body simulation requires all-to-all communication
- There are different ways to implement the communication
 - can use MPI collective communication: broadcast / scatter / gather / all-to-all
 - can use point-to-point communication to implement data exchange between all processes
- Using **collective communication** leads to shorter code and (perhaps) a clearer program structure
 - requires a good understanding of how collective communication works
- Using **point-to-point communication** may be more familiar and is based on simple communication primitives
 - can use same ideas as in the ring communication exercise
 - processes are arranged in a ring where the local bodies are circulated among the processes

Communication structure: ring of processes

- All-to-all communication implemented in a ring of processes

- The positions of the bodies are circulated through a ring

- a process P_i updates the positions of its own local bodies and sends them to the next process in the ring
- P_i calculates the forces between its own local bodies
- when P_i receives the positions from some other process P_j , it calculates the force these cause on its own bodies
- after $P-1$ communication steps, all bodies have visited all processes and the total forces have been calculated
- then the processes can update new velocities for the next timestep



- Communication can proceed simultaneously between all processes
 - all processes have the same amount of work, both computation and communication

Algorithm structure

■ Algorithm for a ring-based solution


- initialize all velocities at time $0.5 \cdot \text{deltat}$
- for each time interval
 - update positions of local bodies
 - calculate forces between the own local bodies
 - for all other processes
 - send local bodies to next process in the ring
 - receive bodies from previous process in the ring
 - compute forces between local bodies and received bodies
 - update velocities of local bodies

■ The force values are accumulated with contributions of particles both from local bodies and from the other processes

- all these are added together to get the total forces on the local particles

Implementation with collective communication

- We can also use MPI collective communication to exchange information about the bodies
 - the communication does not need to do the data exchange in $P-1$ steps, but can directly send the needed data to all processes
 - the implementation is otherwise similar, but the communication structure can be simpler



- The exchange of particle positions is a **gather-operation**
 - coordinates of the local bodies in each process are gathered into one array
 - the result needs to be distributed to *all* processes
- Can implement the data exchange with an *MPI_Gather* followed by a *MPI_Broadcast* of the result
 - or we can use an *MPI_Allgather*

MPI_Gather and MPI_Allgather

- MPI_Gather collects the result into *recvbuf* on the *root*-process
 - `int MPI_Gather(const void *sendbuf, int sendcount, MPI_Datatype sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm)`
- MPI_Allgather has no root-process, the results is in *recvbuf* in all processes
 - `int MPI_Allgather(const void *sendbuf, int sendcount, MPI_Datatype sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype, MPI_Comm comm)`
- Can use *MPI_IN_PLACE* as the value of *sendbuf* on the *root*-process
 - the gather-operation is done in-place, i.e. *recvbuf* is also used as *sendbuf*
 - values of *sendcount* and *sendtype* are ignored, the receive-parameters are used to also determine send-parameters
- When *MPI_IN_PLACE* is used, *sendcount* can be given as 0 and *sendtype* as *MPI_DATATYPE_NULL*

Testing the implementation

- Test the parallel implementation with the same input data as in the sequential program
 - if you don't change the random number generator or the seed values, the same initial positions and masses will be generated
- The program writes out the final positions of the bodies
 - compare these to the result from the sequential program
 - in a correct parallel implemetation, they should be the same at least with 2 decimals of precision
- You can for instance use the Unix command *diff* to compare two files
 - *diff file.txt file2.txt* compares the two text file and prints out the rows that differ
 - *diff -y file1.txt file2.txt* writes out the files in two columns side by side
 - see the manual page for *diff* (*man diff*)